

## Exploratory Analysis & Cleaning

```
In [73]: #Import required Libraries
import numpy as np # Linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [74]: #Loading the dataset
housing=pd.read_csv('HousingDataSet.csv')
```

```
In [75]: #To control the display of decimals. Printed dataframe was showing a lot of de
pd.set_option('float_format', '{:.2f}'.format)
```

```
In [76]: # Print the first 5 rows to understand the dataset
housing.head()
```

```
Out[76]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors
0	7129300520	20141013T000000	221900.00	3	1.00	1180	5650	1.00
1	6414100192	20141209T000000	538000.00	3	2.25	2570	7242	2.00
2	5631500400	20150225T000000	180000.00	2	1.00	770	10000	1.00
3	2487200875	20141209T000000	604000.00	4	3.00	1960	5000	1.00
4	1954400510	20150218T000000	510000.00	3	2.00	1680	8080	1.00

5 rows × 21 columns

```

In [77]: #Format the date to compute age of the hous and rennovation age
#Drop columns with date form and keep calculated age and rennovation age
# As previously done by Sharma (2021)
d = []
for i in housing['date'].values:
    d.append(i[:4])

housing['date'] = d

# convert everything to same datatype
for i in housing.columns:
    housing[i]=housing[i].astype(float)

#make a new column age of the house
housing['age'] = housing['date'] - housing['yr_built']

#calculate the total years of renovation
housing['renov_age'] = np.abs(housing['yr_renovated'] - housing['yr_built'])
housing['renov_age'] = housing.renov_age.apply(lambda x: x if len(str(int(x)))

#remove unwanted columns like yr_built, date, id
housing.drop(['date', 'yr_built', 'yr_renovated'], axis=1, inplace=True)
housing.head()

```

```

Out[77]:

```

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	vie
0	7129300520.00	221900.00	3.00	1.00	1180.00	5650.00	1.00	0.00	0.0
1	6414100192.00	538000.00	3.00	2.25	2570.00	7242.00	2.00	0.00	0.0
2	5631500400.00	180000.00	2.00	1.00	770.00	10000.00	1.00	0.00	0.0
3	2487200875.00	604000.00	4.00	3.00	1960.00	5000.00	1.00	0.00	0.0
4	1954400510.00	510000.00	3.00	2.00	1680.00	8080.00	1.00	0.00	0.0

```
In [78]: # Check dtypes and null values
housing.info()
# Dataframe contains 21 columns and 21.613 rows
# dtype for price,bathrooms and floors are float64.
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 20 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   id                    21613 non-null  float64
 1   price                 21613 non-null  float64
 2   bedrooms              21613 non-null  float64
 3   bathrooms             21613 non-null  float64
 4   sqft_living           21613 non-null  float64
 5   sqft_lot              21613 non-null  float64
 6   floors                21613 non-null  float64
 7   waterfront            21613 non-null  float64
 8   view                  21613 non-null  float64
 9   condition             21613 non-null  float64
10  grade                 21613 non-null  float64
11  sqft_above            21613 non-null  float64
12  sqft_basement         21613 non-null  float64
13  zipcode               21613 non-null  float64
14  lat                   21613 non-null  float64
15  long                  21613 non-null  float64
16  sqft_living15         21613 non-null  float64
17  sqft_lot15            21613 non-null  float64
18  age                   21613 non-null  float64
19  renov_age             21613 non-null  float64
dtypes: float64(20)
memory usage: 3.3 MB
```

```
In [79]: #Converting floats to integers
housing[['price', 'floors', 'bathrooms']] = housing[['price', 'floors', 'bathrooms']
housing.dtypes
```

```
Out[79]: id                float64
price                int32
bedrooms            float64
bathrooms            int32
sqft_living          float64
sqft_lot             float64
floors               int32
waterfront           float64
view                 float64
condition            float64
grade                float64
sqft_above            float64
sqft_basement         float64
zipcode              float64
lat                  float64
long                 float64
sqft_living15         float64
sqft_lot15            float64
age                  float64
renov_age             float64
dtype: object
```

## Null Values

```
In [80]: #Alternative method to confirm that there is not null values in the entire dat
#There is no null values
housing.isnull().values.any()
```

```
Out[80]: False
```

## Find Duplicates

```
In [81]: # Find duplicates in pandas based on Id Column
duplicated=housing.duplicated(subset=['id'],keep='first')
duplicated.sum() # There are 177 duplicates
```

```
Out[81]: 177
```

```
In [82]: ## Printing duplicated rows
duplicated=housing[housing.duplicated(subset=['id'],keep='first')]
print(duplicated)
```

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	\
94	6021501535.00	700000	3.00	1	1580.00	5000.00	
314	4139480200.00	1400000	4.00	3	4290.00	12103.00	
325	7520000520.00	240500	2.00	1	1240.00	12092.00	
346	3969300030.00	239900	4.00	1	1000.00	7134.00	
372	2231500030.00	530000	4.00	2	2180.00	10754.00	
...	...	...	...	...	...	...	
20181	7853400250.00	645000	4.00	3	2910.00	5260.00	
20613	2724049222.00	220000	2.00	2	1000.00	1092.00	
20670	8564860270.00	502000	4.00	2	2680.00	5539.00	
20780	6300000226.00	380000	4.00	1	1200.00	2171.00	
21581	7853420110.00	625000	3.00	3	2780.00	6000.00	

	floors	waterfront	view	condition	grade	sqft_above	sqft_basement
94	1	0.00	0.00	3.00	8.00	1290.00	290.00
314	1	0.00	3.00	3.00	11.00	2690.00	1600.00
325	1	0.00	0.00	3.00	6.00	960.00	280.00
346	1	0.00	0.00	3.00	6.00	1000.00	0.00
372	1	0.00	0.00	5.00	7.00	1100.00	1000.00

```
In [83]: # Remove outliers. We are keeping the last occurrence based on the assumption that
housing.drop_duplicates(subset=['id'],keep='last',inplace=True)
```

```
In [84]: # Dataset now contains 21.436 rows after removing the duplicates
housing.shape
```

Out[84]: (21436, 20)

## Outliers


```
In [85]: # Statistical Analysis
# Price,bedrooms,bathrooms, sqft_living, sqft_lot,sqft_above, sqft_basement var
housing.describe()
```

```
Out[85]:
```

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	wate
count	21436.00	21436.00	21436.00	21436.00	21436.00	21436.00	21436.00	21436.00
mean	4580765328.18	541649.96	3.37	1.75	2082.70	15135.64	1.45	
std	2876589633.67	367314.93	0.93	0.73	919.15	41538.62	0.55	
min	1000102.00	75000.00	0.00	0.00	290.00	520.00	1.00	
25%	2123700078.75	324866.00	3.00	1.00	1430.00	5040.00	1.00	
50%	3904921185.00	450000.00	3.00	2.00	1920.00	7614.00	1.00	
75%	7308675062.50	645000.00	4.00	2.00	2550.00	10696.25	2.00	
max	9900000190.00	7700000.00	33.00	8.00	13540.00	1651359.00	3.00	

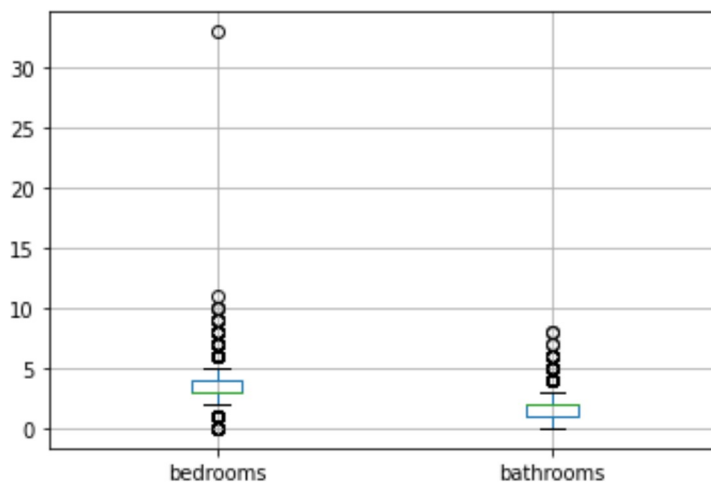
```
In [86]: # box plot price  
housing.boxplot(column='price')
```

Out[86]: <AxesSubplot:>




```
In [87]: # box plots bedrooms and bathrooms  
housing.boxplot(column=['bedrooms', 'bathrooms'])
```

Out[87]: <AxesSubplot:>



```
In [88]: # box plots sqft_living, sqft_above and sqft_basement
housing.boxplot(column=['sqft_living', 'sqft_above', 'sqft_basement'])
```

Out[88]: <AxesSubplot:>



## REMOVING OUTLIERS

The following code is to remove outliers of the dataframe using the IQR alternative

<https://www.youtube.com/watch?v=Vc4cXIAa69Y> (<https://www.youtube.com/watch?v=Vc4cXIAa69Y>)

```
In [89]: ## define a function called outliers which returns a list of index of outliers
def outliers(df, ft):
    Q1 = df[ft].quantile(0.25)
    Q3 = df[ft].quantile(0.75)
    IQR = Q3 - Q1

    lower_bound=Q1-1.5*IQR
    upper_bound=Q3 + 1.5*IQR

    ls=df.index[ (df[ft]<lower_bound)|(df[ft]>upper_bound)]

    return ls
```

```
In [90]: # create an empty list to store the output indices from multiple rows
index_list=[]
for feature in ['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'sqft
```

```
In [91]: index_list
```

```
Out[91]: [5,  
          21,  
          49,  
          69,  
          125,  
          153,  
          216,  
          246,  
          269,  
          270,  
          282,  
          300,  
          312,  
          314,  
          384,  
          419,  
          427,  
          450,  
          472,  
          473]
```

```
In [92]: # The number of items in the list, which represents the number of outliers found  
len(index_list)
```

```
Out[92]: 6148
```

```
In [93]: #define a function called "remove" which returns a cleaned dataframe without outliers  
def remove(df,ls):  
    ls=sorted(set(ls))  
    df= df.drop(ls)  
    return df
```

```
In [94]: # applying the "remove" function created above.  
new_house=remove(housing,index_list)
```



In [95]: `new_house.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 17357 entries, 0 to 21612
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    17357 non-null  float64
1   price                 17357 non-null  int32
2   bedrooms              17357 non-null  float64
3   bathrooms              17357 non-null  int32
4   sqft_living            17357 non-null  float64
5   sqft_lot               17357 non-null  float64
6   floors                 17357 non-null  int32
7   waterfront             17357 non-null  float64
8   view                  17357 non-null  float64
9   condition              17357 non-null  float64
10  grade                  17357 non-null  float64
11  sqft_above              17357 non-null  float64
12  sqft_basement           17357 non-null  float64
13  zipcode                17357 non-null  float64
14  ...                    17357 non-null  float64
```

In [96]: `new_house.describe()`

Out[96]:

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfr
<b>count</b>	17357.00	17357.00	17357.00	17357.00	17357.00	17357.00	17357.00	17357
<b>mean</b>	4755449747.46	460796.25	3.28	1.64	1869.30	7204.29	1.43	C
<b>std</b>	2869137858.73	196947.87	0.78	0.61	660.07	3521.80	0.56	C
<b>min</b>	2800031.00	78000.00	2.00	0.00	440.00	520.00	1.00	C
<b>25%</b>	2321300325.00	308900.00	3.00	1.00	1370.00	4800.00	1.00	C
<b>50%</b>	4068300280.00	425000.00	3.00	2.00	1788.00	7140.00	1.00	C
<b>75%</b>	7507500015.00	577500.00	4.00	2.00	2300.00	9176.00	2.00	C
<b>max</b>	9900000190.00	1125000.00	5.00	3.00	4220.00	19177.00	3.00	1


In [97]: *#Remove column id since it is not relevant for the following algorithms*  
*#Remove sqft\_living, sqft\_lot since there is a more updated version sqft\_livin*  
`new_house.drop(['id', 'sqft_living', 'sqft_lot'], axis=1, inplace=True)`  
`new_house.head()`

Out[97]:

	price	bedrooms	bathrooms	floors	waterfront	view	condition	grade	sqft_above	sqft_ba
<b>0</b>	221900	3.00	1	1	0.00	0.00	3.00	7.00	1180.00	
<b>1</b>	538000	3.00	2	2	0.00	0.00	3.00	7.00	2170.00	
<b>2</b>	180000	2.00	1	1	0.00	0.00	3.00	6.00	770.00	
<b>3</b>	604000	4.00	3	1	0.00	0.00	5.00	7.00	1050.00	
<b>4</b>	510000	3.00	2	1	0.00	0.00	3.00	8.00	1680.00	

# Correlation

```
In [98]: plt.figure(figsize=(20,15))
sns.heatmap(housing.corr(), annot=True)
plt.title("Heatmap of all the Features", fontsize = 25);
plt.show()
```



```
In [99]: #Finding out current data set dimensions
new_house.shape
```

```
Out[99]: (17357, 17)
```

```
In [100]: #Finding correlation coefficient among variables
corr_features = []

for i, r in new_house.corr().iterrows():
    k=0
    for j in range(len(r)):
        if i!= r.index[k]:
            if r.values[k] >=0.5:
                corr_features.append([i, r.index[k], r.values[k]])
        k += 1
corr_features
```

```
Out[100]: [['price', 'grade', 0.5849647065295798],
['price', 'sqft_living15', 0.5040975834167322],
['bathrooms', 'floors', 0.509518412623005],
['bathrooms', 'grade', 0.5088272418632837],
['bathrooms', 'sqft_above', 0.5320829257269674],
['floors', 'bathrooms', 0.509518412623005],
['floors', 'sqft_above', 0.5312115290185615],
['grade', 'price', 0.5849647065295798],
['grade', 'bathrooms', 0.5088272418632837],
['grade', 'sqft_above', 0.6754818964898076],
['grade', 'sqft_living15', 0.6404299564673358],
['sqft_above', 'bathrooms', 0.5320829257269674],
['sqft_above', 'floors', 0.5312115290185615],
['sqft_above', 'grade', 0.6754818964898076],
['sqft_above', 'sqft_living15', 0.7098650293884988],
['sqft_living15', 'price', 0.5040975834167322],
['sqft_living15', 'grade', 0.6404299564673358],
['sqft_living15', 'sqft_above', 0.7098650293884988]]
```

```
In [101]: #Removing highly correlated variables with a coefficient above 0.8 and printing
feat = []
for i in corr_features:
    if i[2] >= 0.8:
        feat.append(i[0])
        feat.append(i[1])

new_house.drop(list(set(feat)), axis=1, inplace=True)
new_house.shape
#Dimensions remain the same beacuse there are zero highly correlated variables
```

```
Out[101]: (17357, 17)
```

```
In [103]: # Export cleaned dataset as a csv
new_house.to_csv(r'C:\Users\may93\Downloads\FinalNew_house.csv', index=False)
```

```
In [ ]:
```

