

Multiple Regression Algorithm

In [81]: *#read excel file*

```
import math
import numpy as np
import pandas as pd
df = pd.read_csv("FinalNew_house.csv")
```

In [82]: *#Load libraries for partition and LinearRegression*

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
from dmbs import regressionSummary, classificationSummary, exhaustive_search,
```

Categorical Variables

In [94]: *# Converting to categorical variables zip code and waterfront*

In [95]: *#Converting categorical variables into integers to remove decimals and then in*

```
df[["zipcode", "waterfront"]] = df[["zipcode", "waterfront"]].astype("int")
df[["zipcode", "waterfront"]] = df[["zipcode", "waterfront"]].astype("category")
```

1. Multiple Regression

In [96]: *# Split the data set into predictors (X) and outcome (y - price)*

```
predictors = df.drop(['price'], axis=1)
outcome = df['price']
```

In [97]: *# Partition data into predictors (x) and output (y)*

```
X = pd.get_dummies(predictors, drop_first=True)
y = outcome
```

In [98]: *#Show first 9 rows of predictors*
X.head(9)

Out[98]:

| | bedrooms | bathrooms | floors | view | condition | grade | sqft_above | sqft_basement | lat | |
|---|----------|-----------|--------|------|-----------|-------|------------|---------------|---------|----|
| 0 | 3.0 | 1 | 1 | 0.0 | 3.0 | 7.0 | 1180.0 | 0.0 | 47.5112 | -1 |
| 1 | 3.0 | 2 | 2 | 0.0 | 3.0 | 7.0 | 2170.0 | 400.0 | 47.7210 | -1 |
| 2 | 2.0 | 1 | 1 | 0.0 | 3.0 | 6.0 | 770.0 | 0.0 | 47.7379 | -1 |
| 3 | 4.0 | 3 | 1 | 0.0 | 5.0 | 7.0 | 1050.0 | 910.0 | 47.5208 | -1 |
| 4 | 3.0 | 2 | 1 | 0.0 | 3.0 | 8.0 | 1680.0 | 0.0 | 47.6168 | -1 |
| 5 | 3.0 | 2 | 2 | 0.0 | 3.0 | 7.0 | 1715.0 | 0.0 | 47.3097 | -1 |
| 6 | 3.0 | 1 | 1 | 0.0 | 3.0 | 7.0 | 1060.0 | 0.0 | 47.4095 | -1 |
| 7 | 3.0 | 1 | 1 | 0.0 | 3.0 | 7.0 | 1050.0 | 730.0 | 47.5123 | -1 |
| 8 | 3.0 | 2 | 2 | 0.0 | 3.0 | 7.0 | 1890.0 | 0.0 | 47.3684 | -1 |

9 rows × 84 columns

In [99]: X.columns

Out[99]: Index(['bedrooms', 'bathrooms', 'floors', 'view', 'condition', 'grade',
'sqft_above', 'sqft_basement', 'lat', 'long', 'sqft_living15',
'sqft_lot15', 'age', 'renov_age', 'waterfront_1', 'zipcode_98002',
'zipcode_98003', 'zipcode_98004', 'zipcode_98005', 'zipcode_98006',
'zipcode_98007', 'zipcode_98008', 'zipcode_98010', 'zipcode_98011',
'zipcode_98014', 'zipcode_98019', 'zipcode_98022', 'zipcode_98023',
'zipcode_98024', 'zipcode_98027', 'zipcode_98028', 'zipcode_98029',
'zipcode_98030', 'zipcode_98031', 'zipcode_98032', 'zipcode_98033',
'zipcode_98034', 'zipcode_98038', 'zipcode_98039', 'zipcode_98040',
'zipcode_98042', 'zipcode_98045', 'zipcode_98052', 'zipcode_98053',
'zipcode_98055', 'zipcode_98056', 'zipcode_98058', 'zipcode_98059',
'zipcode_98065', 'zipcode_98070', 'zipcode_98072', 'zipcode_98074',
'zipcode_98075', 'zipcode_98077', 'zipcode_98092', 'zipcode_98102',
'zipcode_98103', 'zipcode_98105', 'zipcode_98106', 'zipcode_98107',
'zipcode_98108', 'zipcode_98109', 'zipcode_98112', 'zipcode_98115',
'zipcode_98116', 'zipcode_98117', 'zipcode_98118', 'zipcode_98119',
'zipcode_98122', 'zipcode_98125', 'zipcode_98126', 'zipcode_98133',
'zipcode_98136', 'zipcode_98144', 'zipcode_98146', 'zipcode_98148',
'zipcode_98155', 'zipcode_98166', 'zipcode_98168', 'zipcode_98177',
'zipcode_98178', 'zipcode_98188', 'zipcode_98198', 'zipcode_98199'],
dtype='object')

```
In [100]: #Show firs 9 rows of outcomes  
y.describe()
```

```
Out[100]: count      1.735700e+04  
mean        4.607963e+05  
std         1.969479e+05  
min         7.800000e+04  
25%         3.089000e+05  
50%         4.250000e+05  
75%         5.775000e+05  
max         1.125000e+06  
Name: price, dtype: float64
```

```
In [101]: # Split the data into training and validation  
train_X, valid_X, train_y, valid_y = train_test_split(X, y, test_size=0.4, ran
```

```
In [102]: # Built the Linear Model based on the training data  
df_lm = LinearRegression()  
df_lm.fit(train_X, train_y)
```

```
Out[102]: LinearRegression()
```

```
In [92]: #Print coefficients  
print(pd.DataFrame({'Predictor': X.columns, 'coefficient': df_lm.coef_}))
```

| | Predictor | coefficient |
|----|---------------|---------------|
| 0 | bedrooms | -1588.796078 |
| 1 | bathrooms | 10115.824223 |
| 2 | floors | -17319.404817 |
| 3 | view | 31304.483824 |
| 4 | condition | 25286.742243 |
| .. | ... | ... |
| 79 | zipcode_98177 | 208899.773112 |
| 80 | zipcode_98178 | 58407.593922 |
| 81 | zipcode_98188 | 37836.003180 |
| 82 | zipcode_98198 | 29431.368776 |
| 83 | zipcode_98199 | 359451.938714 |

[84 rows x 2 columns]

1. Performance Check

```
In [103]: #Print performance measures(training data)
regressionSummary(train_y, df_lm.predict(train_X))
```

Regression statistics

```

                Mean Error (ME) : -0.0000
      Root Mean Squared Error (RMSE) : 79132.5305
                Mean Absolute Error (MAE) : 58397.2642
                Mean Percentage Error (MPE) : -2.1369
Mean Absolute Percentage Error (MAPE) : 14.0287
```

```
In [109]: #Calculating the R2 score for the model (training data)
df_lm_predt = df_lm.predict(train_X)
score_lr_valt = adjusted_r2_score(train_y, df_lm_predt, df_lm)*100
print("The R2 score for training data is:", score_lr_valt.round(1), "%")
```

The R2 score for training data is: 83.8 %

```
In [106]: #Print performace measures for the validation set
regressionSummary(valid_y, df_lm.predict(valid_X))
```

Regression statistics

```

                Mean Error (ME) : -259.8028
      Root Mean Squared Error (RMSE) : 79984.8699
                Mean Absolute Error (MAE) : 59014.6641
                Mean Percentage Error (MPE) : -2.3362
Mean Absolute Percentage Error (MAPE) : 14.4164
```

```
In [110]: #Calculating the R2 score for the model for the validation set
df_lm_pred = df_lm.predict(valid_X)
score_lr_val = adjusted_r2_score(valid_y, df_lm_pred, df_lm)*100
print("The R2 score for validation data is:", score_lr_val.round(1), "%")
```

The R2 score for validation data is: 83.2 %

```
In [111]: #Make prediction of 2 houses in the validation dataset
df_lm_pred = df_lm.predict(valid_X)
result = pd.DataFrame({'Predicted': df_lm_pred, 'Actual': valid_y, 'Residual':
result.sample(2)
```

```
Out[111]:
```

| | Predicted | Actual | Residual |
|--------------|---------------|--------|--------------|
| 13585 | 397720.602425 | 389999 | -7721.602425 |
| 2103 | 536464.502529 | 560000 | 23535.497471 |

Dimension Reduction - Forward Selection

In [113]: *#Forward Selection*

```

def train_model(variables):
    if len(variables) == 0:
        return None
    model = LinearRegression()
    model.fit(train_X[list(variables)], train_y)
    return model

def score_model(model, variables):
    if len(variables) == 0:
        return AIC_score(train_y, [train_y.mean()] * len(train_y), model, df=1)
    return AIC_score(train_y, model.predict(train_X[variables]), model)

best_model, best_variables = forward_selection(train_X.columns, train_model, s
print("The best variables are:", best_variables)

```

Variables: bedrooms, bathrooms, floors, view, condition, grade, sqft_above, sqft_basement, lat, long, sqft_living15, sqft_lot15, age, renov_age, waterfront_1, zipcode_98002, zipcode_98003, zipcode_98004, zipcode_98005, zipcode_98006, zipcode_98007, zipcode_98008, zipcode_98010, zipcode_98011, zipcode_98014, zipcode_98019, zipcode_98022, zipcode_98023, zipcode_98024, zipcode_98027, zipcode_98028, zipcode_98029, zipcode_98030, zipcode_98031, zipcode_98032, zipcode_98033, zipcode_98034, zipcode_98038, zipcode_98039, zipcode_98040, zipcode_98042, zipcode_98045, zipcode_98052, zipcode_98053, zipcode_98055, zipcode_98056, zipcode_98058, zipcode_98059, zipcode_98065, zipcode_98070, zipcode_98072, zipcode_98074, zipcode_98075, zipcode_98077, zipcode_98092, zipcode_98102, zipcode_98103, zipcode_98105, zipcode_98106, zipcode_98107, zipcode_98108, zipcode_98109, zipcode_98112, zipcode_98115, zipcode_98116, zipcode_98117, zipcode_98118, zipcode_98119, zipcode_98122, zipcode_98125, zipcode_98126, zipcode_98133, zipcode_98136, zipcode_98144, zipcode_98146, zipcode_98148, zipcode_98155, zipcode_98166, zipcode_98168, zipcode_98177, zipcode_98178, zipcode_98188, zipcode_98198, zipcode_98199

Start: score=283512.01, constant

Step: score=279063.19, add grade

Step: score=275775.39, add lat

Step: score=274534.09, add age

Step: score=273425.89, add sqft_living15

Step: score=272938.20, add bathrooms

Step: score=272516.19, add zipcode_98004

Step: score=272210.92, add zipcode_98155

Step: score=271873.59, add zipcode_98133

Step: score=271546.93, add sqft_above

Step: score=271078.70, add sqft_basement

Step: score=270740.74, add zipcode_98040

Step: score=270393.49, add view

Step: score=270062.64, add zipcode_98028

Step: score=269776.05, add zipcode_98034

Step: score=269504.69, add zipcode_98112

Step: score=269246.33, add zipcode_98019

Step: score=268961.20, add zipcode_98125

Step: score=268664.00, add zipcode_98011

Step: score=268329.06, add zipcode_98072

Step: score=268049.80, add zipcode_98177

Step: score=267765.08, add condition

Step: score=267625.11, add zipcode_98119

```
Step: score=267478.93, add zipcode_98109
Step: score=267344.85, add zipcode_98022
Step: score=267208.86, add zipcode_98116
Step: score=267078.14, add zipcode_98102
Step: score=266942.31, add zipcode_98006
Step: score=266815.53, add zipcode_98199
Step: score=266702.00, add zipcode_98178
Step: score=266597.96, add zipcode_98168
Step: score=266489.00, add zipcode_98058
Step: score=266396.46, add zipcode_98077
Step: score=266307.45, add zipcode_98014
Step: score=266219.94, add zipcode_98136
Step: score=266135.73, add zipcode_98122
Step: score=266051.25, add zipcode_98010
Step: score=265970.80, add zipcode_98105
Step: score=265892.28, add zipcode_98005
Step: score=265813.07, add zipcode_98027
Step: score=265746.62, add zipcode_98108
Step: score=265687.22, add zipcode_98055
Step: score=265624.76, add zipcode_98106
Step: score=265567.35, add waterfront_1
Step: score=265520.24, add zipcode_98029
Step: score=265454.06, add long
Step: score=265412.67, add zipcode_98188
Step: score=265372.52, add zipcode_98033
Step: score=265332.87, add zipcode_98031
Step: score=265294.14, add zipcode_98056
Step: score=265257.81, add zipcode_98039
Step: score=265227.79, add renov_age
Step: score=265197.39, add zipcode_98146
Step: score=265171.10, add zipcode_98052
Step: score=265145.12, add zipcode_98059
Step: score=265125.06, add zipcode_98038
Step: score=265106.94, add floors
Step: score=265087.17, add zipcode_98198
Step: score=265066.71, add zipcode_98032
Step: score=265047.69, add zipcode_98118
Step: score=265030.04, add zipcode_98053
Step: score=265016.83, add zipcode_98030
Step: score=264999.86, add zipcode_98042
Step: score=264985.62, add zipcode_98074
Step: score=264970.56, add zipcode_98065
Step: score=264961.47, add zipcode_98148
Step: score=264950.55, add zipcode_98023
Step: score=264941.04, add zipcode_98166
Step: score=264932.66, add zipcode_98092
Step: score=264925.37, add zipcode_98024
Step: score=264920.13, add zipcode_98070
Step: score=264913.84, add zipcode_98003
Step: score=264908.86, add zipcode_98126
Step: score=264906.06, add zipcode_98002
Step: score=264905.19, add zipcode_98144
Step: score=264903.76, add zipcode_98103
Step: score=264901.26, add zipcode_98115
Step: score=264895.68, add zipcode_98107
Step: score=264870.10, add zipcode_98117
```

```

Step: score=264834.04, add zipcode_98045
Step: score=264814.32, add zipcode_98007
Step: score=264787.26, add zipcode_98008
Step: score=264640.48, add zipcode_98075
Step: score=264640.48, add None
The best variables are: ['grade', 'lat', 'age', 'sqft_living15', 'bathrooms',
'zipcode_98004', 'zipcode_98155', 'zipcode_98133', 'sqft_above', 'sqft_baseme
nt', 'zipcode_98040', 'view', 'zipcode_98028', 'zipcode_98034', 'zipcode_9811
2', 'zipcode_98019', 'zipcode_98125', 'zipcode_98011', 'zipcode_98072', 'zipc
ode_98177', 'condition', 'zipcode_98119', 'zipcode_98109', 'zipcode_98022', '
zipcode_98116', 'zipcode_98102', 'zipcode_98006', 'zipcode_98199', 'zipcode_9
8178', 'zipcode_98168', 'zipcode_98058', 'zipcode_98077', 'zipcode_98014', 'z
ipcode_98136', 'zipcode_98122', 'zipcode_98010', 'zipcode_98105', 'zipcode_98
005', 'zipcode_98027', 'zipcode_98108', 'zipcode_98055', 'zipcode_98106', 'wa
terfront_1', 'zipcode_98029', 'long', 'zipcode_98188', 'zipcode_98033', 'zipc
ode_98031', 'zipcode_98056', 'zipcode_98039', 'renov_age', 'zipcode_98146', '
zipcode_98052', 'zipcode_98059', 'zipcode_98038', 'floors', 'zipcode_98198',
'zipcode_98032', 'zipcode_98118', 'zipcode_98053', 'zipcode_98030', 'zipcode_
98042', 'zipcode_98074', 'zipcode_98065', 'zipcode_98148', 'zipcode_98023', '
zipcode_98166', 'zipcode_98092', 'zipcode_98024', 'zipcode_98070', 'zipcode_9
8003', 'zipcode_98126', 'zipcode_98002', 'zipcode_98144', 'zipcode_98103', 'z
ipcode_98115', 'zipcode_98107', 'zipcode_98117', 'zipcode_98045', 'zipcode_98
007', 'zipcode_98008', 'zipcode_98075']

```

2. Multiple Regression (only variables)

```

In [114]: # Use only the first 1000 records and select columns for regression analysis
predictors2 = df[['grade', 'lat', 'age', 'sqft_living15']]
outcome2 = df['price']

```

```

In [115]: # Partition data into predictors (x) and output (y)
X2 = predictors2
y2 = outcome

```

```

In [116]: # Split the data into training and validation
train_X2, valid_X2, train_y2, valid_y2 = train_test_split(X2, y2, test_size=0.

```

```

In [117]: # Built the Linear Model based on the training data
df_lm2 = LinearRegression()
df_lm2.fit(train_X2, train_y2)

```

Out[117]: LinearRegression()

```

In [118]: #Print coefficients
print(pd.DataFrame({'Predictor': X2.columns, 'coefficient': df_lm2.coef_}))

```

| | Predictor | coefficient |
|---|---------------|---------------|
| 0 | grade | 109584.085093 |
| 1 | lat | 534014.865066 |
| 2 | age | 1894.846443 |
| 3 | sqft_living15 | 94.982097 |

2. Performance Check

```
In [48]: #Print performance measures(training data)
         regressionSummary(train_y2, df_lm2.predict(train_X2))
```

Regression statistics

```

                        Mean Error (ME) : -0.0000
          Root Mean Squared Error (RMSE) : 122563.9553
                Mean Absolute Error (MAE) : 93327.2997
                Mean Percentage Error (MPE) : -6.6471
Mean Absolute Percentage Error (MAPE) : 22.6443
```

```
In [120]: #Calculating the R2 score for the model (training data)
          df_lm_predt2 = df_lm2.predict(train_X2)
          score_lr_valt2 = adjusted_r2_score(train_y2, df_lm_predt2, df_lm2)*100
          print("The R2 score for training data is:", score_lr_valt2.round(1), "%")
```

The R2 score for training data is: 62.0 %

```
In [121]: #Print performace measures for the validation set
          regressionSummary(valid_y2, df_lm2.predict(valid_X2))
```

Regression statistics

```

                        Mean Error (ME) : -39.6458
          Root Mean Squared Error (RMSE) : 122659.9450
                Mean Absolute Error (MAE) : 92964.2270
                Mean Percentage Error (MPE) : -6.5821
Mean Absolute Percentage Error (MAPE) : 22.5140
```

```
In [122]: #Calculating the R2 score for the model
          df_lm_pred2 = df_lm2.predict(valid_X2)
          score_lr_val2 = adjusted_r2_score(valid_y2, df_lm_pred2, df_lm2)*100
          print("The R2 score for validation data is:", score_lr_val2.round(1), "%")
```

The R2 score for validation data is: 60.9 %

```
In [123]: #Make prediction of 2 houses in the validation dataset
          df_lm_pred2 = df_lm2.predict(valid_X2)
          result2 = pd.DataFrame({'Predicted': df_lm_pred2, 'Actual': valid_y2, 'Residual':
                                result2.sample(2)})
```

```
Out[123]:
```

| | Predicted | Actual | Residual |
|-------|---------------|--------|---------------|
| 15476 | 202577.965503 | 205000 | 2422.034497 |
| 1371 | 412948.231053 | 338000 | -74948.231053 |

In []:

